

## Team Member

Alan Xing ID: 101144347

Chien-I Chao ID: 101135366

Nicholas Rasmussen ID: 1839739

GitHub Link : <https://github.com/Sup3000gt/Machine-Learning-Project>

## Convolutional Neural Network Project

```
In [58]: # import all libraries

from ucimlrepo import fetch_ucirepo
import numpy as np
from sklearn.model_selection import train_test_split

from tensorflow.keras import Input, Model
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, BatchNormalization, Activation, MaxPool
from sklearn.model_selection import KFold
from sklearn.metrics import accuracy_score, recall_score, precision_score, f1_score

import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix
```

### Step 1.1 Loading the Dataset

```
In [59]: # Using the code from ucimlrepo directly load the data

optical_recognition_of_handwritten_digits = fetch_ucirepo(id=80)

# data (as pandas dataframes)
X = optical_recognition_of_handwritten_digits.data.features
y = optical_recognition_of_handwritten_digits.data.targets
```

### Let's look at the data

```
In [60]: print(X.shape)
```

(5620, 64)

```
In [61]: X.head()
```

Out[61]:

	Attribute1	Attribute2	Attribute3	Attribute4	Attribute5	Attribute6	Attribute7	Attrib
0	0	1	6	15	12	1	0	
1	0	0	10	16	6	0	0	
2	0	0	8	15	16	13	0	
3	0	0	0	3	11	16	0	
4	0	0	5	14	4	0	0	

5 rows × 64 columns



In [62]: `y.head()`

Out[62]:

	class
0	0
1	0
2	7
3	4
4	6

## Step 1.2 Preprocessing - Normalization

Here we normalize all values between zero and one, as is standard for neural network processing

In [63]:

```
max_pixel_value = X.max().max()
print("Maximum pixel value:", max_pixel_value)

# seem like the Max value of pixel is 16 instead of 255 like other image, we will n
X_normalized = X / max_pixel_value
```

Maximum pixel value: 16

In [64]:

```
# after Normalization
X_normalized.head()
```

Out[64]:

	Attribute1	Attribute2	Attribute3	Attribute4	Attribute5	Attribute6	Attribute7	Attrib
0	0.0	0.0625	0.3750	0.9375	0.7500	0.0625	0.0	
1	0.0	0.0000	0.6250	1.0000	0.3750	0.0000	0.0	
2	0.0	0.0000	0.5000	0.9375	1.0000	0.8125	0.0	
3	0.0	0.0000	0.0000	0.1875	0.6875	1.0000	0.0	
4	0.0	0.0000	0.3125	0.8750	0.2500	0.0000	0.0	

5 rows × 64 columns



## Step 1.2 Preprocessing - Reshape

Here we add the channels dimension to all of our data points

```
In [65]: """
current shape is 8 x 8 array, (64) we need to reshape it into 2D array
"""
# Data frame do not have reshape, so we will convert into numpy array first
X_np = X_normalized.to_numpy()

X_resaped = X_np.reshape(-1, 8, 8, 1) # Reshape to (5620, 8, 8, 1)

print("Reshaped data shape:", X_resaped.shape)
```

Reshaped data shape: (5620, 8, 8, 1)

```
In [66]: # convert y to numpy array as well
y = y.to_numpy()
```

## Step 1.2 Preprocessing - Data Split

Here we split the data

```
In [67]: # Now we are ready to split the data
X_train, X_test, y_train, y_test = train_test_split(X_resaped, y, test_size=0.2, r
```

## Step 1.3 and 1.4 Data Augmentation and One-Hot Encoding

Here we add Gaussian noise to the training set while also doubling it. Then use a categorical representation for the label data

```
In [68]: # Create a function to add Gaussian noise to an image
def add_gaussian_noise(image, mean=0, std=0.05):
    noise = np.random.normal(mean, std, image.shape)
    noisy_image = image + noise
    return np.clip(noisy_image, 0, 1) # Clip values to [0, 1]

# Create a larger training set by adding noise to each image
```

```

X_train_augmented = []
y_train_augmented = []

for i in range(2):
    for i in range(len(X_train)):
        original_image = X_train[i]
        noisy_image = add_gaussian_noise(original_image)
        X_train_augmented.append(noisy_image)
        y_train_augmented.append(y_train[i])

# Convert lists to numpy arrays
X_train_augmented = np.array(X_train_augmented)
y_train_augmented = np.array(y_train_augmented)

# One-hot encode the Labels
y_train_augmented_onehot = to_categorical(y_train_augmented, num_classes=10)
y_test_onehot = to_categorical(y_test, num_classes=10)

```

## Step 2/3/4 CNN Architecture/ Max Pooling/ Softmax

Here, we create our Convolutional Neural Network. All output dimension sizes have been documented. Parameters are in the model summary before training

```

In [71]: def cnet(input_shape=(8,8,1), num_classes=10, dropout_rate=0.2):
# Define the input Layer
inputs = Input(shape=input_shape) #Outputs a (8,8,1) matrix

# Define the encoder part
x = Conv2D(32, 3, padding="same")(inputs) #Outputs a (8,8,32) matrix
x = BatchNormalization()(x)
x = Activation("relu")(x)
x = MaxPooling2D((2, 2), padding="same")(x) #Outputs a (4,4,32) matrix, halving
x = SpatialDropout2D(dropout_rate)(x) #We also use spatial dropout to remove

x = Conv2D(64, 3, padding="same")(x) #Outputs a (4,4,64) matrix
x = BatchNormalization()(x)
x = Activation("relu")(x)
x = MaxPooling2D((2, 2), padding="same")(x) #Outputs a (2,2,64) matrix
x = Dropout(dropout_rate)(x) #Normal dropout that drop neurons to zero

x = Conv2D(128, 3, padding="same")(x) #Outputs a (2,2,128) matrix
x = BatchNormalization()(x)
x = Activation("relu")(x)
x = Flatten()(x) #Outputs a 512 sized sector from the dimensions 2x2x128
x = Dropout(dropout_rate)(x) #Normal dropout that drop neurons to zero

x = Dense(64)(x) #Outputs a 64 sized sector
x = Dropout(dropout_rate)(x) #Normal dropout that drop neurons to zero

# Define the output Layer
outputs = Dense(num_classes, activation="softmax")(x) #Outputs a vector of size

# Create the model
model = Model(inputs=inputs, outputs=outputs)

```

```
return model
```

```
In [72]: #code found online for reinitializing the weights of a model in a method
def reset_weights(model):
    for layer in model.layers:
        if isinstance(layer, Model):
            reset_weights(layer)
            continue
        for k, initializer in layer.__dict__.items():
            if "initializer" not in k:
                continue
            # find the corresponding variable
            var = getattr(layer, k.replace("_initializer", ""))
            var.assign(initializer(var.shape, var.dtype))

model = cnet()

# Compile our model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

model.summary()
```

Model: "model\_4"

Layer (type)	Output Shape	Param #
=====		
input_5 (InputLayer)	[(None, 8, 8, 1)]	0
conv2d_12 (Conv2D)	(None, 8, 8, 32)	320
batch_normalization_12 (Batch Normalization)	(None, 8, 8, 32)	128
activation_12 (Activation)	(None, 8, 8, 32)	0
max_pooling2d_8 (MaxPooling2D)	(None, 4, 4, 32)	0
spatial_dropout2d_4 (Spatial Dropout)	(None, 4, 4, 32)	0
conv2d_13 (Conv2D)	(None, 4, 4, 64)	18496
batch_normalization_13 (Batch Normalization)	(None, 4, 4, 64)	256
activation_13 (Activation)	(None, 4, 4, 64)	0
max_pooling2d_9 (MaxPooling2D)	(None, 2, 2, 64)	0
dropout_12 (Dropout)	(None, 2, 2, 64)	0
conv2d_14 (Conv2D)	(None, 2, 2, 128)	73856
batch_normalization_14 (Batch Normalization)	(None, 2, 2, 128)	512
activation_14 (Activation)	(None, 2, 2, 128)	0
flatten_4 (Flatten)	(None, 512)	0
dropout_13 (Dropout)	(None, 512)	0
dense_8 (Dense)	(None, 64)	32832
dropout_14 (Dropout)	(None, 64)	0
dense_9 (Dense)	(None, 10)	650
=====		
Total params: 127050 (496.29 KB)		
Trainable params: 126602 (494.54 KB)		
Non-trainable params: 448 (1.75 KB)		

## Step 5. Model Performance - Training

```
In [73]: reset_weights(model)
history = model.fit(
    X_train_augmented, y_train_augmented_onehot,
    epochs=20,
    validation_data=(X_test, y_test_onehot)
)
```

Epoch 1/20

```
c:\users\nick_\appdata\local\programs\python\python38\lib\site-packages\keras\src\initializers\initializers.py:120: UserWarning: The initializer GlorotUniform is unseeded and being called multiple times, which will return identical values each time (even if the initializer is unseeded). Please update your code to provide a seed to the initializer, or avoid using the same initializer instance more than once.
  warnings.warn(
```

281/281 [=====] - 2s 4ms/step - loss: 0.4845 - accuracy: 0.8429 - val\_loss: 0.6839 - val\_accuracy: 0.8657  
Epoch 2/20  
281/281 [=====] - 1s 4ms/step - loss: 0.1448 - accuracy: 0.9513 - val\_loss: 0.1023 - val\_accuracy: 0.9662  
Epoch 3/20  
281/281 [=====] - 1s 4ms/step - loss: 0.0866 - accuracy: 0.9698 - val\_loss: 0.0511 - val\_accuracy: 0.9858  
Epoch 4/20  
281/281 [=====] - 1s 4ms/step - loss: 0.0839 - accuracy: 0.9729 - val\_loss: 0.0463 - val\_accuracy: 0.9893  
Epoch 5/20  
281/281 [=====] - 1s 4ms/step - loss: 0.0618 - accuracy: 0.9793 - val\_loss: 0.0612 - val\_accuracy: 0.9840  
Epoch 6/20  
281/281 [=====] - 1s 4ms/step - loss: 0.0627 - accuracy: 0.9806 - val\_loss: 0.0271 - val\_accuracy: 0.9911  
Epoch 7/20  
281/281 [=====] - 1s 4ms/step - loss: 0.0540 - accuracy: 0.9835 - val\_loss: 0.0411 - val\_accuracy: 0.9902  
Epoch 8/20  
281/281 [=====] - 1s 4ms/step - loss: 0.0362 - accuracy: 0.9877 - val\_loss: 0.0434 - val\_accuracy: 0.9920  
Epoch 9/20  
281/281 [=====] - 1s 4ms/step - loss: 0.0458 - accuracy: 0.9858 - val\_loss: 0.0365 - val\_accuracy: 0.9884  
Epoch 10/20  
281/281 [=====] - 1s 4ms/step - loss: 0.0425 - accuracy: 0.9870 - val\_loss: 0.0357 - val\_accuracy: 0.9893  
Epoch 11/20  
281/281 [=====] - 1s 4ms/step - loss: 0.0289 - accuracy: 0.9903 - val\_loss: 0.0244 - val\_accuracy: 0.9947  
Epoch 12/20  
281/281 [=====] - 1s 4ms/step - loss: 0.0367 - accuracy: 0.9884 - val\_loss: 0.0321 - val\_accuracy: 0.9929  
Epoch 13/20  
281/281 [=====] - 1s 4ms/step - loss: 0.0280 - accuracy: 0.9904 - val\_loss: 0.0184 - val\_accuracy: 0.9929  
Epoch 14/20  
281/281 [=====] - 1s 4ms/step - loss: 0.0198 - accuracy: 0.9932 - val\_loss: 0.0217 - val\_accuracy: 0.9947  
Epoch 15/20  
281/281 [=====] - 1s 4ms/step - loss: 0.0285 - accuracy: 0.9904 - val\_loss: 0.0440 - val\_accuracy: 0.9893  
Epoch 16/20  
281/281 [=====] - 1s 4ms/step - loss: 0.0285 - accuracy: 0.9901 - val\_loss: 0.0228 - val\_accuracy: 0.9938  
Epoch 17/20  
281/281 [=====] - 1s 4ms/step - loss: 0.0251 - accuracy: 0.9923 - val\_loss: 0.0364 - val\_accuracy: 0.9911  
Epoch 18/20  
281/281 [=====] - 1s 4ms/step - loss: 0.0220 - accuracy: 0.9939 - val\_loss: 0.0308 - val\_accuracy: 0.9911  
Epoch 19/20  
281/281 [=====] - 1s 4ms/step - loss: 0.0284 - accuracy: 0.9910 - val\_loss: 0.0284 - val\_accuracy: 0.9947



Epoch 20/20  
281/281 [=====] - 1s 4ms/step - loss: 0.0233 - accuracy: 0.9915 - val\_loss: 0.0290 - val\_accuracy: 0.9938

## Step 5.1 Model Performance with Accuracy plot, Loss plot, and Confusion matrix



- As we can see above, our model performance is close to 99%, let's visualize the results of our model


```
In [74]: # Plotting the training and validation accuracy
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(loc='upper left')
plt.show()

# Plotting the training and validation loss
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(loc='upper left')
plt.show()

predictions = model.predict(X_test)
predicted_classes = np.argmax(predictions, axis=1)
true_classes = np.argmax(to_categorical(y_test, num_classes=10), axis=1)

# Confusion matrix
cm = confusion_matrix(true_classes, predicted_classes)
sns.heatmap(cm, annot=True, fmt='d')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
```





## Step 5.2 K-Fold cross-validation

- Here, we implement K-Fold cross-validation. We also collect pertinent metrics for each class across all folds.

```
In [75]: # Initialize lists to store metrics across folds
accuracy_scores = []
sensitivity_scores = []
specificity_scores = []
f1_scores = []

kf = KFold(n_splits=5)
for train_index, test_index in kf.split(X_resaped):
    reset_weights(model)
    X_train, X_test = X_resaped[train_index], X_resaped[test_index]
    y_train, y_test = y[train_index], y[test_index]
    # Create a larger training set by adding noise to each image
    X_train_augmented = []
    y_train_augmented = []
    y_test_one = to_categorical(y_test, num_classes=10)

    for i in range(2):
        for i in range(len(X_train)):
            original_image = X_train[i]
            noisy_image = add_gaussian_noise(original_image)
            X_train_augmented.append(noisy_image)
            y_train_augmented.append(y_train[i])
```

```

# Convert Lists to numpy arrays
X_train_augmented = np.array(X_train_augmented)
y_train_augmented = np.array(y_train_augmented)
model.fit(X_train_augmented, to_categorical(y_train_augmented, num_classes=10),

# Make predictions on the test set
y_pred = model.predict(X_test)
y_pred_classes = np.argmax(y_pred, axis=1)

# Compute classification report (includes precision, recall, F1-score, and supp
class_report = classification_report(y_test, y_pred_classes, target_names=["Cla

print("Classification Report:")
print(class_report)

# Parse the classification report to extract relevant metrics
lines = class_report.split("\n")
metrics = {}
for line in lines[2:-5]: # Skip header and footer lines
    class_name, precision, recall, f1_score, _ = line.split()
    metrics[class_name] = {
        "Precision": float(precision),
        "Recall": float(recall),
        "F1-score": float(f1_score),
    }
...

# Now you can access the metrics for each class
for class_name, values in metrics.items():
    print(f"Metrics for {class_name}:")
    print(f"Precision: {values['Precision']:.4f}")
    print(f"Recall (Sensitivity): {values['Recall']:.4f}")
    print(f"F1-score: {values['F1-score']:.4f}")
    print()
...

```

c:\users\nick\_\appdata\local\programs\python\python38\lib\site-packages\keras\src\initializers\initializers.py:120: UserWarning: The initializer GlorotUniform is unseeded and being called multiple times, which will return identical values each time (even if the initializer is unseeded). Please update your code to provide a seed to the initializer, or avoid using the same initializer instance more than once.

```
warnings.warn(
```

36/36 [=====] - 0s 2ms/step

Classification Report:

	precision	recall	f1-score	support
Class0	1.00	0.98	0.99	110
Class1	0.98	1.00	0.99	118
Class2	1.00	1.00	1.00	114
Class3	0.98	0.99	0.99	107
Class4	0.99	0.98	0.99	113
Class5	0.99	0.96	0.98	113
Class6	0.99	1.00	1.00	111
Class7	0.99	1.00	1.00	118
Class8	1.00	0.97	0.99	105
Class9	0.96	0.99	0.97	115
accuracy			0.99	1124
macro avg	0.99	0.99	0.99	1124
weighted avg	0.99	0.99	0.99	1124

c:\users\nick\_\appdata\local\programs\python\python38\lib\site-packages\keras\src\initializers\initializers.py:120: UserWarning: The initializer GlorotUniform is unseeded and being called multiple times, which will return identical values each time (even if the initializer is unseeded). Please update your code to provide a seed to the initializer, or avoid using the same initializer instance more than once.

warnings.warn(

36/36 [=====] - 0s 2ms/step

Classification Report:

	precision	recall	f1-score	support
Class0	1.00	0.99	1.00	114
Class1	0.97	0.99	0.98	114
Class2	1.00	1.00	1.00	115
Class3	0.99	0.99	0.99	123
Class4	0.98	1.00	0.99	106
Class5	0.99	1.00	1.00	107
Class6	1.00	0.98	0.99	111
Class7	0.99	1.00	1.00	113
Class8	0.99	0.97	0.98	108
Class9	1.00	0.99	1.00	113
accuracy			0.99	1124
macro avg	0.99	0.99	0.99	1124
weighted avg	0.99	0.99	0.99	1124

c:\users\nick\_\appdata\local\programs\python\python38\lib\site-packages\keras\src\initializers\initializers.py:120: UserWarning: The initializer GlorotUniform is unseeded and being called multiple times, which will return identical values each time (even if the initializer is unseeded). Please update your code to provide a seed to the initializer, or avoid using the same initializer instance more than once.

warnings.warn(

36/36 [=====] - 0s 2ms/step

Classification Report:

	precision	recall	f1-score	support
Class0	1.00	1.00	1.00	104
Class1	1.00	0.98	0.99	106
Class2	0.98	1.00	0.99	106
Class3	0.99	0.98	0.99	109
Class4	0.99	0.98	0.99	124
Class5	0.98	0.99	0.98	120
Class6	0.98	0.99	0.99	114
Class7	0.99	1.00	1.00	113
Class8	1.00	0.99	1.00	114
Class9	0.99	0.98	0.99	114
accuracy			0.99	1124
macro avg	0.99	0.99	0.99	1124
weighted avg	0.99	0.99	0.99	1124

c:\users\nick\_\appdata\local\programs\python\python38\lib\site-packages\keras\src\initializers\initializers.py:120: UserWarning: The initializer GlorotUniform is unseeded and being called multiple times, which will return identical values each time (even if the initializer is unseeded). Please update your code to provide a seed to the initializer, or avoid using the same initializer instance more than once.

warnings.warn(

36/36 [=====] - 0s 2ms/step

Classification Report:

	precision	recall	f1-score	support
Class0	0.99	1.00	1.00	116
Class1	0.98	0.98	0.98	119
Class2	0.99	0.99	0.99	113
Class3	0.99	0.99	0.99	119
Class4	0.99	1.00	1.00	111
Class5	1.00	0.97	0.99	104
Class6	0.99	0.99	0.99	108
Class7	1.00	0.94	0.97	109
Class8	1.00	0.98	0.99	119
Class9	0.91	1.00	0.95	106
accuracy			0.98	1124
macro avg	0.99	0.98	0.98	1124
weighted avg	0.99	0.98	0.98	1124

c:\users\nick\_\appdata\local\programs\python\python38\lib\site-packages\keras\src\initializers\initializers.py:120: UserWarning: The initializer GlorotUniform is unseeded and being called multiple times, which will return identical values each time (even if the initializer is unseeded). Please update your code to provide a seed to the initializer, or avoid using the same initializer instance more than once.

warnings.warn(

36/36 [=====] - 0s 2ms/step

Classification Report:

	precision	recall	f1-score	support
Class0	1.00	1.00	1.00	110
Class1	0.97	1.00	0.98	114
Class2	1.00	1.00	1.00	109
Class3	0.97	1.00	0.98	114
Class4	1.00	0.99	1.00	114
Class5	0.98	0.99	0.99	114
Class6	0.98	1.00	0.99	114
Class7	0.99	1.00	1.00	113
Class8	1.00	0.94	0.97	108
Class9	0.99	0.96	0.97	114
accuracy			0.99	1124
macro avg	0.99	0.99	0.99	1124
weighted avg	0.99	0.99	0.99	1124

## Step 5.3 Metrics Analysis Across Folds

Here, we collect the average, range, and standard deviation of all the metrics we collect.

```
In [76]: # Calculate range and standard deviation for precision, recall, and F1-score
precision_values = [values["Precision"] for values in metrics.values()]
recall_values = [values["Recall"] for values in metrics.values()]
f1_values = [values["F1-score"] for values in metrics.values()]

avg_precision = np.mean(precision_values)
avg_recall = np.mean(recall_values)
avg_f1 = np.mean(f1_values)

range_precision = np.ptp(precision_values)
range_recall = np.ptp(recall_values)
range_f1 = np.ptp(f1_values)

std_precision = np.std(precision_values)
std_recall = np.std(recall_values)
std_f1 = np.std(f1_values)

print("\nOverall Metrics Across Each Fold for all Classes:")

print(f"Average Precision: {avg_precision:.4f}")
print(f"Range of Precision: {range_precision:.4f}")
print(f"Standard Deviation of Precision: {std_precision:.4f}")

print(f"Average Recall (Sensitivity): {avg_recall:.4f}")
print(f"Range of Recall: {range_recall:.4f}")
print(f"Standard Deviation of Recall: {std_recall:.4f}")

print(f"Average F1-score: {avg_f1:.4f}")
print(f"Range of F1-score: {range_f1:.4f}")
print(f"Standard Deviation of F1-score: {std_f1:.4f}")
```

Overall Metrics Across Each Fold for all Classes:

Average Precision: 0.9880

Range of Precision: 0.0300

Standard Deviation of Precision: 0.0117

Average Recall (Sensitivity): 0.9880

Range of Recall: 0.0600

Standard Deviation of Recall: 0.0199

Average F1-score: 0.9880

Range of F1-score: 0.0300

Standard Deviation of F1-score: 0.0117

In [ ]: